

Hashing

David Croft

Coventry University

david.croft@coventry.ac.uk

2016

1 Finding things

2 Hashes

- What is a hash
- Load
- Collisions
- Using them in code

3 Bloom filters



Already seen binary search faster than linear.

- But binary search only works on ordered sequences.
 - Sorted list/array, Binary Search Trees (BSTs) etc.
- $O(\log n)$ vs. $O(n)$.
- What's better than $O(n)$?
 - $O(1)$.

Can we lookup values in $O(1)$ time?

Unordered sequence

C

Want to store information on all UK motorways.

- 49 motorways.

Option 1.

- Unordered sequence.
- List/array/vector.
- Finding specific motorway is $O(n)$.
- Space required, $O(n)$.

Pos	Motorways
0	M9
1	M55
2	M898
3	M4
4	M1
5	M6
...	
45	M2
46	M56
47	M53
48	M3

Ordered sequence

C

Option 2.

- Ordered sequence.
- Sorted List/array/vector or BST.
- Finding specific motorway is $O(\log n)$.
- Space required, $O(n)$.

Pos	Motorways
0	M1
1	M2
2	M3
3	M4
4	M5
5	M6
...	
45	M606
46	M621
47	M876
48	M898

Lookup table

C

Option 3

- Lookup table.
- Each motorway stored in position corresponding to it's number.
 - E.g. M1 in position 1, M53 in position 53.
- Finding specific motorway is $O(1)$.
 - Very fast.
- Space required, $O(\max(n))$, 899 spaces.
 - Very inefficient...in this case.
 - Can be VERY efficient, massive time savings for small memory cost.

Pos	Motorways
0	
1	M1
2	M2
3	M3
...	
53	M53
54	M54
55	M55
56	M56
...	
894	
895	
896	
897	
898	M898

Option 4

- Hash table.
- Pass each motorway through a hash function.
- Store hashes in lookup table.
- Finding specific motorway is $O(1)$ ish.
- Space required, $O(n)$.

Hash tables.

Unordered associative arrays.

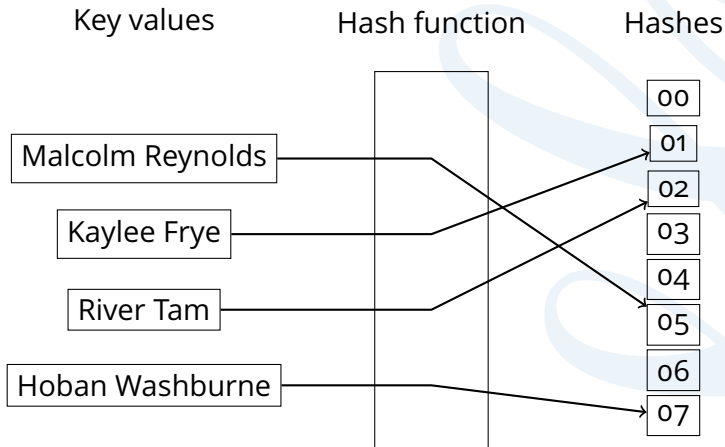
- Unordered - we have no control over the item orders.
- Associative - Lookup a value based on a key.
 - I.e. Python `dict()`, C++ `map<>`.
- Fast.
 - $O(1)$ lookup (potentially).
 - $O(1) \leq \text{Reality} \leq O(n)$.

Hash is just a number.

- Based on some other value.
- Motorway example hash could just be the M-number.
 - I.e. M898 \rightarrow 898.


What if key is not an int?

- Hash function.
- Converts an input of any size/range to a fixed size/range.
- Related to, but distinct from:
 - Checksums.
 - Fingerprints.
 - Parity codes.



Signs of a good hash table hashing algorithms.

- Computationally lightweight.
- Evenly distributed hashes.
 - E.g. `len()` would be terrible hash function, loads of different inputs produce same value.

Input	CRC32 hash
"Small text."	3840495446
1	3523407757
[1,2,3,4,5]	1191942644
"On the Origin of Species" 158,454 words	3877468994
	192271774

So how does a hash function work?

- Depends on hash function and purpose.
- Not going to be implementing any real algorithms.
 - Optimized to be very fast, not understandable.
 - Generally full of binary representations.
 - Bit shifting.
- Simple hashing algorithm (division method).
 - 1 Break the thing being hashed into blocks.
 - 1, 2, 4, 8, 16 bytes in size.
 - 2 Add up all the blocks.
 - 3 Modulo by a prime number.

Cryptographic hashes.

A

Side note, different hashes for different purposes.

- Hash table hashes.
 - Computationally lightweight as possible.
- Cryptographic hashes.
 - Computationally lightweight-ish to go key \rightarrow hash.
 - Computationally expensive to go hash \rightarrow key.
 - MD (Message-Digest algorithm)
 - Famously MD5, widely used, no longer secure.
 - MD6 still good.
 - SHA (Secure Hash Algorithm)
 - SHA-0, SHA-1 not secure.
 - SHA-2, SHA-3 still good.

Wait a second

1

Hang on a minute, there's problem.

- Hashed "buckeroo" with CRC32.
 - CRC32 **NOT** best choice for hash table but is easy and widespread.
 - Hash of 1306201125
 - So our hash table needs at least 1,306,201,125 slots.
 - 4 bytes per integer * 1306201125 slots = 5.2 gigabytes.
 - Not going to work.
 - Solution? Take the modulo of the hash.
 - Create table of small size.
 - `slot = hash % len(hashtable)`

Converting big sequences into short hashes.

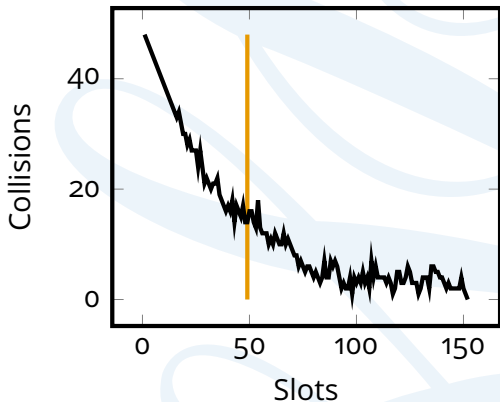
- Any downsides?
- Some distinct sequences **MUST** produce same hash.
- Hash collision.
- I.e. Hashing an `int`
 - `int` has 4 billion possible values.
 - If hash is one byte, then 256 possible slots.
 - 4 billion possible values in 256 possible slots, collisions will happen.
- E.g. CRC32 hash.
 - "plumless" → 1306201125 ← "buckeroo"

Bigger the hash table == less collisions.

- Ideal motorway hash table.
 - 49 motorways.
 - 49 slots.
 - 0 collisions.
- Reality.
 - 49 slots.
 - 14 collisions.
 - I.e. M18, M67 and M606 in slot 18.

Table size

1



- Optimal.
 - 152 slots.
 - 0 collisions.
 - Size only 3.1 times ideal.

Number of collisions depends on load factor (l).

- Ratio of elements (n) to available slots (k).

- $l = \frac{n}{k}$

- High load = lots of collisions (probably).

- Low load = few collisions (probably).

- $l > 1.0$ = definitely some collisions.

- Previous motorway example.

- 49 motorways, 49 slots, $l = 1$

- 49 motorways, 152 slots, $l = 0.31$

Real world hash tables.

- Automatically resize to provide more slots as load increases.
- Advantages
 - Table size adjusts for the amount of elements stored in it.
 - Minimal wasted memory.
- Disadvantages
 - Have to shuffle everything around when the table resizes.
 - Not all that time consuming.

So what do we do when we have collisions?

- If have ≥ 1 elements then collisions are possible.
 - Regardless of table size.
- Two main approaches

Separate chaining.

- Each slot is a linked list.
- Infinitely resizeable.
- Add new item to end of the list.

Open addressing.

- Slot is already full?
- Try next slot until an empty one is found.

Handle collisions II

I

Separate chaining.

buckeroo 1306201125

 \Rightarrow

Slot	
0	→ limpet →
1	→ zombie →
2	→
3	→ plumless → buckeroo
4	→
5	→ gondola →

Open addressing

buckeroo 1306201125

 \Rightarrow

Slot	
0	limpet
1	zombie
2	
3	plumless
4	buckeroo
5	gondala

Dictionaries!

Hash tables in Python

C

```
import sys

def main():
    motorways = {}

    motorways["M1"] = (193.5, 1959)
    motorways["M2"] = (25.7, 1963)
    motorways["M3"] = (58.6, 1971)
    motorways["M4"] = (191.9, 1961)
    # [...]
    motorways["M898"] = (0.5, 1985)

    print( 'The %s is %0.1f miles long' % ("M4",
↪ motorways["M4"][0]) )
    print( 'The %s opened in %d' % ("M898",
↪ motorways["M898"][1]) )

if __name__ == '__main__':
    sys.exit(main())
```

Maps!

```
#include <iostream>
#include <map>
using namespace std;

int main()
{
    map< string, pair<float,int> > motorways;

    motorways.emplace( "M1", make_pair<float,int>(193.5, 1959) );
    //                      [...]
    motorways.emplace( "M898", make_pair<float,int>( 0.5, 1985 ) );

    cout << "The " << "M1" << " is " <<
        motorways.find("M1")->second.first << " miles long" << endl;

    cout << "The " << "M898" << " opened in " <<
        motorways.find("M898")->second.second << endl;

    return 0;
}
```

Hashes have many other possible applications.

- Finding duplicates.
 - Hash table but count number of things in each slots.
- Similarity comparisons.
 - E.g. Soundex, Metaphone.
 - Names that sound the same have same hash.
- Image recognition.
- Bloom filters.
 - Are almost magic.

Harry Potter
and the
Bloom Filter

Neat trick with hashes.


- Can 'store' 1000 things in the space for 100.
 - Doesn't actually store the items.
 - Can say if an element is not a member of a set.
 - Can say if element is probably a member of a set

To add a value to the filter.

- Hash the value using multiple different functions.
- Mark the slots for each of those hashes.

To test for a value in the filter.

- Hash the value using all the functions.
 - If not all the slots are marked then value not in filter.
 - If all slots are marked then value *probably* in filter.



Bloom filters are almost magic.

- 1000 items
- $\leq 1\%$ error
- 7 hash functions
- 9586 slots
 - 1 bit per slot.
- 'Store' 1000 integers in the space for 300.
 - Bigger variables mean bigger savings.

Hashing

David Croft

Finding
things

Hashes

What is a hash

Load

Collisions

Using them in code

Bloom filters

Quiz

Hashing

David Croft

Finding
things

Hashes

What is a hash

Load

Collisions

Using them in code

Bloom filters



Recap

The end of 122COM... ...or is it?